

Throughput mean data transfer speed ----(copy past/ Downloads speed)

EBC

SSD HDD

In AWS, **EBS Volume Types** are categorized based on performance and cost, allowing users to select storage optimized for their workload requirements. Below is a detailed explanation of each volume type you listed:

1. General Purpose SSD (gp2):

- **Purpose:** Balances price and performance for a wide variety of workloads.
- **Performance:**
 - Provides baseline performance of 3 IOPS per GB.
 - Can burst up to 3,000 IOPS for volumes under 1 TB.
- **Use Cases:**
 - System boot volumes.
 - Low-latency interactive applications.
 - Development and testing environments.
- **Capacity:** 1 GiB to 16 TiB.
- **Throughput:** Up to 250 MiB/s.

2. General Purpose SSD (gp3):

- **Purpose:** Next-generation general-purpose SSD offering predictable performance and lower cost compared to gp2.
- **Performance:**
 - Delivers a baseline performance of 3,000 IOPS and 125 MiB/s.
 - Performance can be provisioned up to 16,000 IOPS and 1,000 MiB/s, independent of volume size.
- **Use Cases:**
 - Databases, boot volumes, and applications requiring consistent performance.
- **Capacity:** 1 GiB to 16 TiB.
- **Throughput:** Up to 1,000 MiB/s.
- **Key Benefit:** More cost-effective and flexible compared to gp2.

3. Provisioned IOPS SSD (io1):

- **Purpose:** High-performance SSD for mission-critical workloads that require low-latency and high IOPS.
- **Performance:**
 - Provision IOPS from 100 to 64,000 (max 50 IOPS per GiB for volumes up to 64 TiB).
 - Designed for latency-sensitive workloads.
- **Use Cases:**
 - Large-scale databases (e.g., Oracle, SQL Server, MySQL).
 - Applications requiring high transactional workloads.
- **Capacity:** 4 GiB to 16 TiB.
- **Key Benefit:** Customizable performance for demanding workloads.

4. Provisioned IOPS SSD (io2):

- **Purpose:** Enhanced version of io1 with higher durability and better performance.
- **Performance:**
 - Provision up to 64,000 IOPS (max 500 IOPS per GiB for volumes up to 64 TiB).
 - 99.999% durability (compared to io1's 99.9%).
- **Use Cases:**
 - Critical applications requiring very high durability and performance.
 - Enterprise-grade databases and ERP systems.
- **Capacity:** 4 GiB to 64 TiB.
- **Key Benefit:** Higher durability and efficiency for high-performance workloads.

5. Cold HDD (sc1):

- **Purpose:** Low-cost, infrequent access storage for data that is rarely accessed.
- **Performance:**
 - Baseline throughput of 12 MiB/s per TiB.
 - Max throughput of 250 MiB/s.
- **Use Cases:**

- Cold data storage.
- Log storage or backups that are rarely accessed.
- **Capacity:** 125 GiB to 16 TiB.
- **Key Benefit:** Cost-efficient storage for rarely accessed data.

6. Throughput Optimized HDD (st1):

- **Purpose:** Optimized for workloads requiring high throughput rather than IOPS.
- **Performance:**
 - Baseline throughput of 40 MiB/s per TiB.
 - Max throughput of 500 MiB/s.
- **Use Cases:**
 - Big data processing (e.g., Hadoop).
 - Data warehouses.
 - Streaming workloads requiring high throughput.
- **Capacity:** 125 GiB to 16 TiB.
- **Key Benefit:** Cost-effective for high-throughput workloads.

7. Magnetic (Standard):

- **Purpose:** Legacy storage option for workloads requiring cost-effective, low-performance storage.
- **Performance:**
 - Average IOPS of 100.
 - Suitable for sequential read/write operations.
- **Use Cases:**
 - Archive storage.
 - Non-critical workloads or boot volumes for older instance types.
- **Capacity:** Up to 1 TiB.
- **Key Benefit:** Simple and cost-effective, but mostly replaced by newer options.

Comparison Summary:

Volume Type	Max IOPS	Max Throughput	Durability	Cost	Use Case
General Purpose SSD (gp2)	16,000	250 MiB/s	Standard	Moderate	Balanced workloads
General Purpose SSD (gp3)	16,000	1,000 MiB/s	Standard	Low	Databases, boot volumes
Provisioned IOPS SSD (io1)	64,000	1,000 MiB/s	High	High	Critical databases
Provisioned IOPS SSD (io2)	64,000	1,000 MiB/s	Very High	High	Enterprise applications
Cold HDD (sc1)	~500	250 MiB/s	Standard	Very Low	Rarely accessed data
Throughput Optimized HDD (st1)	~500	500 MiB/s	Standard	Low	Streaming, big data
Magnetic (Standard)	~100	Low	Standard	Very Low	Legacy, basic storage

Notes:

- **IOPS (Input/Output Operations Per Second):** Measures the performance of a volume in handling read/write operations.
- **Durability:** Higher durability reduces the risk of data loss.
- **Throughput:** Indicates the data transfer rate, often critical for workloads like streaming or big data.
- **Cost Optimization:** Match the volume type to your workload to avoid unnecessary costs.

Let me know if you need further clarification!

Step of how to Attach volume EBS

- Create a instance Launch the instance
- Go to EBS Volume
- Create the Volume
 - Select the Volume types
 - General Purpose SSD (gp2)
 - General Purpose SSD (gp3)
 - Provisioned IOPS SSD (io1)
 - Provisioned IOPS SSD (io2) ----- (input output per sec.)
 - Cold HDD (Sc1)
 - Throughput Optimized HDD (st1)
 - Magnetic (standard)
 - Select Size (GiB)
 - Select IOPS
 - Select throughput -----(data transfer speed MB per sec.)

- Select Availability Zone
- Click the click the Create Volume
- After create the Volume click the Actions
- and select the Attach Volume
 - Provide the default Volume ID
 - Default Availability Zone
 - Select the Running Instance show available in Availability Zone
 - Provide the device name
 - Click the Attach Volume
- Go to the terminal

Command

- lsblk
 - display information about block devices. Block devices include hard drives, SSDs, USB drives, and partitions
 - Syntax `lsblk [option]`
 - Option
 - `lsblk -a` ----- (-a: show all devices, including empty ones)
 - `lsblk -l` ----- (-l: display output in list format)
 - `lsblk -f` ----- (-f: show the file system type UUIDs)

Example of output

```
root@ip-172-31-6-109:/# lsblk
NAME        MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
loop0         7:0    0  26.3M 1 loop /snap/amazon-ssm-agent/9881
loop1         7:1    0  73.9M 1 loop /snap/core22/1663
loop2         7:2    0  38.8M 1 loop /snap/snapd/21759
xvda         202:0    0    8G  0 disk
├─xvda1       202:1    0    7G  0 part /
├─xvda14      202:14   0    4M  0 part
├─xvda15      202:15   0   106M 0 part /boot/efi
└─xvda16      259:0    0   913M 0 part /boot
xvdbd        202:14080 0   160G 0 disk
```

Explanation of Output:

- **NAME:** The name of the block device (e.g., xvda, xvdf).
 - **MAJ:MIN:** Major and minor device numbers.
 - **RM:** Indicates whether the device is removable (1 for removable, 0 otherwise).
 - **SIZE:** The size of the device or partition.
 - **RO:** Indicates whether the device is read-only (1 for read-only, 0 otherwise).
 - **TYPE:** The type of device (e.g., disk, part for partition, rom for CD-ROM).
 - **MOUNTPOINT:** The directory where the device is mounted.
- `mkfs.ext4 /dev/xvdbd`
 - is used to format a block device (in this case, /dev/xvdbd) with the **ext4** file system on a Linux system, such as one running on an AWS EC2 instance

Breakdown of the Command:

- **mkfs.ext4:** Creates an ext4 file system on the specified device.
- **/dev/xvdbd:** Specifies the block device to be formatted. This might represent an attached EBS volume or another storage device.

```
root@ip-172-31-6-109:/home/ubuntu# mkfs.ext4 /dev/xvdbd
mke2fs 1.47.0 (5-Feb-2023)
/dev/xvdbd contains a ext4 file system
created on Thu Jan  2 07:17:42 2025
Proceed anyway? (y,N) y
Creating filesystem with 41943040 4k blocks and 10485760 inodes
Filesystem UUID: 8981afee-3ce6-4a1d-ad75-5d070f0e744e
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632, 2654208,
    4096000, 7962624, 11239424, 20480000, 23887872

Allocating group tables: done
Writing inode tables: done
Creating journal (262144 blocks): done
Writing superblocks and filesystem accounting information: done
```

- `mkdir /test`
- `mount /dev/xvdbd /test`

```

root@ip-172-31-6-109:/# lsblk
NAME        MAJ:MIN     RM  SIZE RO TYPE MOUNTPOINTS
loop0        7:0         0 26.3M  1 loop /snap/amazon-ssm-agent/9881
loop1        7:1         0 73.9M  1 loop /snap/core22/1663
loop2        7:2         0 38.8M  1 loop /snap/snapd/21759
xvda         202:0       0    8G  0 disk
├─xvda1      202:1       0    7G  0 part /
├─xvda14     202:14      0    4M  0 part
├─xvda15     202:15      0 106M  0 part /boot/efi
└─xvda16     259:0       0  913M  0 part /boot
xvdbd        202:14080   0 160G  0 disk /test
root@ip-172-31-6-109:/#

```

- mountpoint /test
 - The EBS volume are allot to the test directory

```

root@ip-172-31-6-109:/# mountpoint /test/
/test/ is a mountpoint
root@ip-172-31-6-109:/#

```

- umount /test
 - Remove the mount

```

root@ip-172-31-6-109:/# umount /test/
root@ip-172-31-6-109:/# lsblk
NAME        MAJ:MIN     RM  SIZE RO TYPE MOUNTPOINTS
loop0        7:0         0 26.3M  1 loop /snap/amazon-ssm-agent/9881
loop1        7:1         0 73.9M  1 loop /snap/core22/1663
loop2        7:2         0 38.8M  1 loop /snap/snapd/21759
xvda         202:0       0    8G  0 disk
├─xvda1      202:1       0    7G  0 part /
├─xvda14     202:14      0    4M  0 part
├─xvda15     202:15      0 106M  0 part /boot/efi
└─xvda16     259:0       0  913M  0 part /boot
xvdbd        202:14080   0 160G  0 disk
root@ip-172-31-6-109:/#

```

- file -s /dev/xvdbd

- how to increase the size of EBS volume
 - Select the volume click the actions
 - Click Modify volume
 - Change the size increase in GB
 - And click modify
 - After increase the volume
 - Go to the terminal
 - Write the command
 - resize2fs /dev/xvdbd ---()
 - df -h --- (Use df -h to confirm the file system has been resized)
 - fdisk /dev/xvdbd
 - n ----(new partition)
 - p ----(primary partition)
 - w ----(save)
 - q ----(quit)
 - mkfs.ext4 dev/xvdbd1 ----(make a file system)
 - mkdir ebs
 - mount /dev/xvdbd1 ebs
 - df -h
 - mountpoint
 - umount -----(remove the mount)